

W 8 - 12 lecture

W 8 - 11 work

W 8, 10, 12 applied

> W8 GPU vs NPU

> W9 Synchronisation, Mutex, deadlocks

> W10 Election Algorithms, Concurrency control

> W11 Matrix Multiplication, Partitioning

> W12 Exponential growth vs Performance Planning

end of today

12/10/24

notes, workshop q,  
summary formulas

> GPU vs NPU vs CPU

CPU  
General tasks.  
↳ Many Instructions

GPU  
Dumb. lots  
of repetitive tasks

NPU  
↳ AI tasks

Lower bandwidth

Very high bandwidth

ILP

TLP

7W9 : Synchronisation, Mutual exclusion, deadlocks

Synchronisation → why? We need to sync processes for many things. To order events, bug fix etc.

How? 2 main types:

logical clocks:

keep ordering of local events on a single node, by relative ordering of events.

physical clocks:

Sync events to a physical clock. UTC

More simple  
Partial ordering

VS

More complicated to sync b/w  
processes  
Total ordering

Algorithms:

Lamport's Algorithm

↳ determine event order only

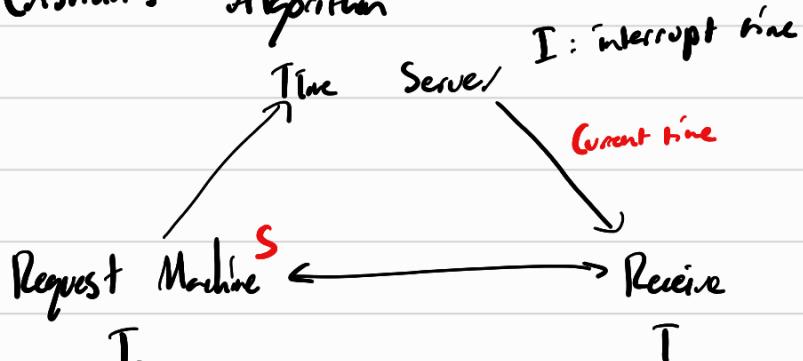
A → B A happens before B

↳ when:

- A + B on same process, A occurs before B
- A sends message, B receives
- If X, Y are independent on diff processors → concurrent  $X \not\rightarrow Y, Y \not\rightarrow X$

if  $A \rightarrow B$ , then  $C(A) \subset C(B)$   
but not other way

Grisian's Algorithm



$$S_{time} = curr\_time + \frac{T_1 - T_0 - I}{2}$$

Averaging Algorithm

- Resync every fixed interval  $T_R$   
↳ set local time using average
- $P_0$  shall be advanced by  $\Delta t_s$

$$\Delta t_s = \frac{t_0 + t_1 + t_2 + t_3}{4} - t_0$$

Berkley's Algorithm  $\rightarrow$  Master polls slaves periodically, sends what slave has to adjust by

- Find time difference b/w master and slave processes.
- Average

$$\text{Corrected time} = \bar{T}_m + \text{Average time}$$

$$\Delta \bar{T}_n = \text{Corrected time} - T_n$$

$\Rightarrow$  WQ

Mutual exclusion: A concurrency control property  
 $\downarrow$   
to prevent race conditions

Ensure concurrent accesses of processes to a shared resource is serialised.

$\hookrightarrow$  Only one process allowed to execute in a critical section

$\hookrightarrow$  segment of code executed by multiple concurrent threads/processes accessing shared resources

Ways of ensuring Mutex:

$\Rightarrow$  Centralised Algorithm:

$\hookrightarrow$  A process is a 'co-ordinator'

1) If process enters critical section, sends request to co-ordinator

2) If no other process, send grant

3) Otherwise put in queue

4) Once process exits, process sends release

Pros: - First-come, first serve (fair)

- Simple + easy

- Only **3 messages** per critical section use: request, queue, release

Cons: - If co-ordinator fails, system may fail

- Process cannot distinguish b/w crashed 10 or 0' permission denied

- Single co-ord performance bottleneck

## → Distributed Algorithm

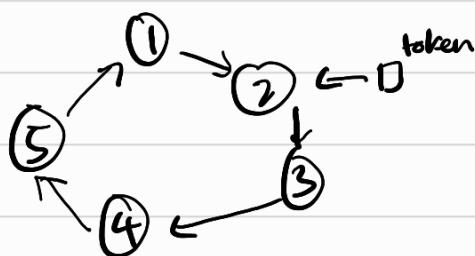
- ↳ Process sends message to all other process,  
containing: critical section name + Process no. + curr time
- ↳ All other processes receive
  - if process not in critical section  
↳ send Ok
  - else check process msg send time vs received msg time  
if rec time < send time  
↳ Ok
  - else → put at end of queue

**if one process does not respond, system doesn't continue**

**$2(n-1)$  messages**       $n = \text{no. processes}$

## Ring-Token:

- Nodes in ring. One token held and passed b/w processes.
- ↳ if node has token and needs critical section, it can
- ↳ otherwise pass token in ring



## → Summary

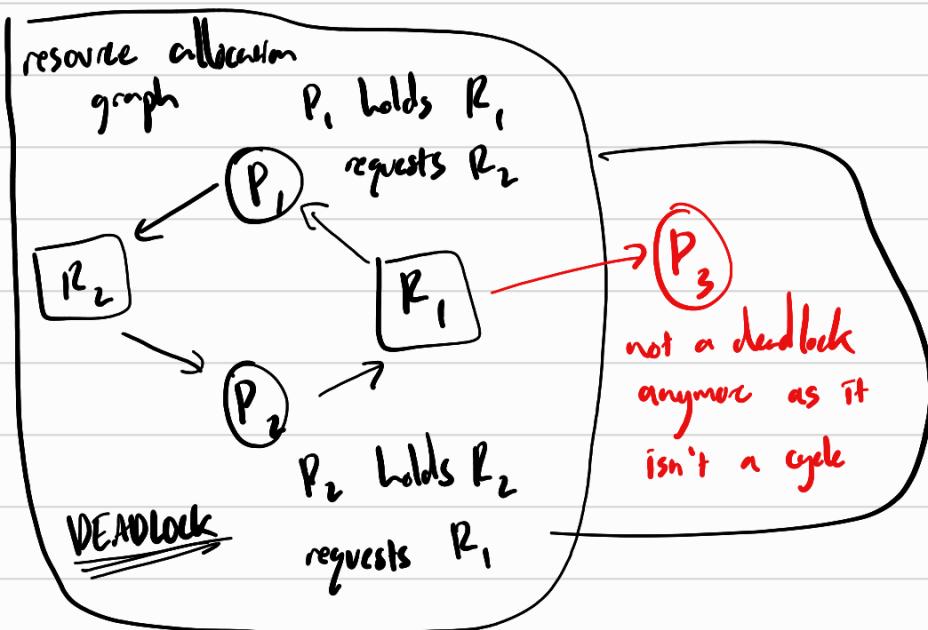
	Messages	Problems
Centralised	3	-Co-ordinator crash / bottleneck
Decentralised	$3n k$ ( $k=1, 2 \dots$ )	-Starvation / low efficiency
Distributed	$2(n-1)$	-Crash of any process
Token-ring	1 to $\infty$	-Lose token / process crash

→ Deadlocks : When a process cannot proceed because it needs resource held by other process which current process holds.

### Conditions:

- 1) Mutual Exclusion: Resource held by at most 1 process
- 2) Hold and Wait: Process that holds resource is waiting for another
- 3) Non-preemption: A resource can only be released by process after completion
- 4) Circular wait:  $P_0 \rightarrow P_1 \rightarrow \dots \rightarrow P_0$

$P_0$  waits for  $P_1$ , waits for  $P_2$  ... waits for  $P_0$



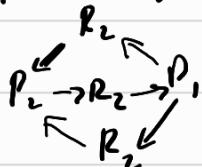
### Deadlocks in graph

if 1 copy of all resources:

↳ cycle

if multiple copies:

↳ knot : All vertices in knot have outgoing edge that terminate at another vertex in knot.



Wait for graph: remove resources



## → Handling Deadlocks

- Ostrich algorithm: ignore problem

→ Detection: Usually easier than prevention

↳ after: kill processes or abort transactions

### - Centralised detection

↳ Each machine maintains resource graph

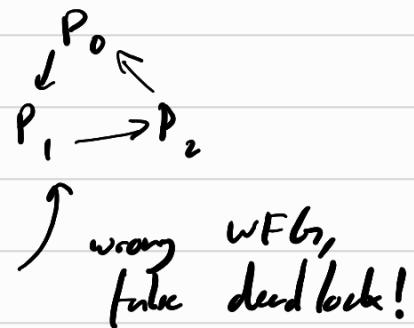
↳ Co-ordinator merges graphs of machines together

↳ Message sent to co-ordinator to update

### - False deadlocks:

If  $P_1$  releases  $P_0$  first,  
then  $P_1$  asks  $P_2$ .

↳ Message sent can detect  
 $P_1$  waits  $P_2$  before  $P_1$  releases  $P_0$

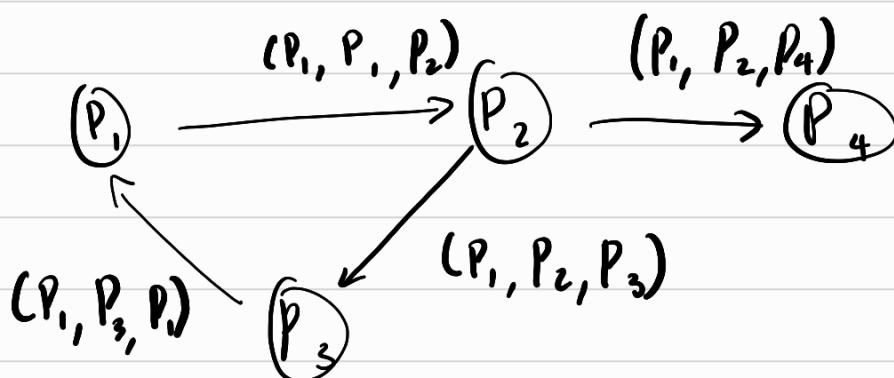


↳ Must use global ordering  
Coordinator or ask each machine if it has release

### - Distributed - CNTL

Probe message generated when process needs to wait for resource.

↳ Contains:  
- Process being blocked  
- Process sending message  
- Process to send to



→ W 10

Election Algorithms: To decide a coordinator  
- responsible for ensuring mutex and detecting deadlocks  
in centralised algorithms

Assumptions:

- Processes have unique id.
- All processes know each other
- Process do not know which other process is alive.

Election:

- 1) Maximum number process selected
- 2) Elected to be coordinator
- 3) All processes must agree

→ Bully Algorithm (Ivanica Molina)

- ↳ 1) Process P sends election msg to higher processes
- 2) If higher process responds → new leader, ok message sent to P
- 3) No response → P wins election  
*(timeout)* ↳ send coordinator message

*(cons):*

- Timeout overhead - slow
- May have delayed response after timeout

→ Ring algorithm

↳ ring topology of processes.

- 1) Process finds communicator is dead
- 2) Send election msg
- 3) Passed down until it reaches original sender
- 4) Coordinator msg passed through signalling new coordinator

Dead processes → Process needs to wait for timeout before bypassing process,  
↳ delays

→ Concurrency Control

↳ Ensure multiple transactions can execute without data inconsistencies / interference.

→ Conflicts

Read-Write Conflict: Shared - exclusive conflict

Write-Write: exclusive - exclusive

Methods to manage conflicts:

→ Lock based

↳ Read locks: Data only to be read

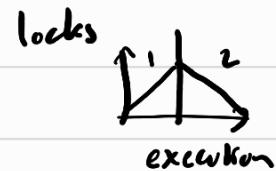
Shared Mode → can be accessed by other

↳ Write locks: Data only to be updated

Exclusive Mode → cannot be accessed until lock removed

2-phase locking: GUARANTEES Serializability

1) Gains all locks needed



2) Releases locks as it goes

→ Timestamp: Uses ordering (logical or physical)

T can read x when:

$$T \geq x_{\text{write time}}$$

T can write to x:

$$T \geq x_{\text{write time}}$$

and

$$T \geq x_{\text{read time}}$$

→ Optimistic: Assumes no conflicts until end

3 phases: 1) read

check no conflicts

Validate: checks for any  $T_i, T_j$  start either 2) validate

1)  $T_i$  finish 3 phase before  $T_j$  or

3) write → to database

2)  $T_i$  finish before  $T_j$  write or 3)  $T_i$  finish read before  $T_j$  read

